

# Machine learning Group8 - Lab1 block 2


*Group-8*

*November 28, 2017*

## Assignment 1

### Adaboost

#### Error Rate vs Tree Level



Loss Function for the selected family is  $LossFunction = (y - f)^2$

## Random Forest




## Assignment 2

### Mixture Model

For K=2


```
## iteration: 1 log likelihood: -1710.416
## iteration: 2 log likelihood: -1708.212
## iteration: 3 log likelihood: -1705.908
## iteration: 4 log likelihood: -1688.444
## iteration: 5 log likelihood: -1568.968
## iteration: 6 log likelihood: -1067.461
## iteration: 7 log likelihood: -294.244
## iteration: 8 log likelihood: 58.06127
## iteration: 9 log likelihood: 109.9385
## iteration: 10 log likelihood: 114.3999
## iteration: 11 log likelihood: 114.3286
```



For K=3

```
## iteration: 1 log likelihood: -2116.967
## iteration: 2 log likelihood: -2113.26
## iteration: 3 log likelihood: -2109.279
## iteration: 4 log likelihood: -2078.903
## iteration: 5 log likelihood: -1879.558
## iteration: 6 log likelihood: -1156.973
## iteration: 7 log likelihood: -299.6505
## iteration: 8 log likelihood: 3.388555
## iteration: 9 log likelihood: 64.81808
## iteration: 10 log likelihood: 81.13098
## iteration: 11 log likelihood: 88.18021
## iteration: 12 log likelihood: 92.315
## iteration: 13 log likelihood: 95.11267
## iteration: 14 log likelihood: 97.15403
## iteration: 15 log likelihood: 98.71629
## iteration: 16 log likelihood: 99.95156
## iteration: 17 log likelihood: 100.9507
## iteration: 18 log likelihood: 101.7718
## iteration: 19 log likelihood: 102.4538
## iteration: 20 log likelihood: 103.0246
## iteration: 21 log likelihood: 103.5045
## iteration: 22 log likelihood: 103.9095
## iteration: 23 log likelihood: 104.2518
## iteration: 24 log likelihood: 104.5414
## iteration: 25 log likelihood: 104.7866
## iteration: 26 log likelihood: 104.994
## iteration: 27 log likelihood: 105.1692
## iteration: 28 log likelihood: 105.3169
## iteration: 29 log likelihood: 105.4411
```

```
## iteration: 30 log likelihood: 105.5451
## iteration: 31 log likelihood: 105.6317
```




For K=4

```
## iteration: 1 log likelihood: -2404.83
## iteration: 2 log likelihood: -2400.939
## iteration: 3 log likelihood: -2396.925
## iteration: 4 log likelihood: -2365.575
## iteration: 5 log likelihood: -2149.829
## iteration: 6 log likelihood: -1317.208
## iteration: 7 log likelihood: -375.1612
## iteration: 8 log likelihood: -64.09552
## iteration: 9 log likelihood: 6.910748
## iteration: 10 log likelihood: 30.05542
## iteration: 11 log likelihood: 40.76704
## iteration: 12 log likelihood: 46.91076
## iteration: 13 log likelihood: 51.04236
## iteration: 14 log likelihood: 54.18383
## iteration: 15 log likelihood: 56.78241
## iteration: 16 log likelihood: 59.0457
## iteration: 17 log likelihood: 61.07583
## iteration: 18 log likelihood: 62.92663
## iteration: 19 log likelihood: 64.62901
## iteration: 20 log likelihood: 66.20267
## iteration: 21 log likelihood: 67.66166
## iteration: 22 log likelihood: 69.01716
## iteration: 23 log likelihood: 70.27908
## iteration: 24 log likelihood: 71.45698
## iteration: 25 log likelihood: 72.56097
## iteration: 26 log likelihood: 73.60253
## iteration: 27 log likelihood: 74.59534
```

```

## iteration: 28 log likelihood: 75.55626
## iteration: 29 log likelihood: 76.50611
## iteration: 30 log likelihood: 77.46999
## iteration: 31 log likelihood: 78.47646
## iteration: 32 log likelihood: 79.5548
## iteration: 33 log likelihood: 80.72966
## iteration: 34 log likelihood: 82.01349
## iteration: 35 log likelihood: 83.39881
## iteration: 36 log likelihood: 84.85428
## iteration: 37 log likelihood: 86.32793
## iteration: 38 log likelihood: 87.75797
## iteration: 39 log likelihood: 89.08701
## iteration: 40 log likelihood: 90.27337
## iteration: 41 log likelihood: 91.29589
## iteration: 42 log likelihood: 92.15209
## iteration: 43 log likelihood: 92.85274
## iteration: 44 log likelihood: 93.4158
## iteration: 45 log likelihood: 93.86161
## iteration: 46 log likelihood: 94.2098
## iteration: 47 log likelihood: 94.47779
## iteration: 48 log likelihood: 94.68027
## iteration: 49 log likelihood: 94.82922
## iteration: 50 log likelihood: 94.93427
## iteration: 51 log likelihood: 95.00306

```



Ideally, mixture model should run for parameters that are not too small nor too big to avoid the underfitting and overfitting of the model. For too few parameters that is for  $K=2$ , the loglikelihood function runs for 11 iterations giving  $\mu$  near to the true values of  $\mu$  while for too many parameters ( $K=4$ ) the convergence steps increases resulting in overfitting. For  $K=3$ , the loglikelihood value increases significantly giving  $\mu_2$  that is  $\mu$  for  $K=2$  and  $\mu_3$  that is  $\mu$  for  $K=3$  near the true values.

## APPENDIX

```
## Question 1

spambase <- read.csv2("spambase.csv", header = TRUE, sep = ";", quote = "\"",
                      dec = ",", fill = TRUE)
spambase <- as.data.frame(spambase)

### Adaboost

n=dim(spambase)[1]
set.seed(12345)
id=sample(1:n, floor(n*2/3))
train=spambase[id,]
test=spambase[-id,]

number_of_trees <- seq(from = 10,to = 100, by = 10)

adaboost <- function(ntrees)
{
  fit <- blackboost(Spam ~., data = train,
                    control = boost_control(mstop = ntrees, nu=0.1),
                    family = Gaussian())

  ypredict <- predict(fit, test)

  error <- mean((test$Spam - ypredict)^2)
}

error_rates_a <- sapply(number_of_trees, adaboost)

plot(error_rates_a,type = "b",main="Error Rate vs Tree Level", xlab= "Number of Trees",
      ylab= "Error",ylim=c(0.05,0.15), col="red", pch=19, cex=1)

### Random Forest

training = sample(1:n,floor(n*2/3))

random_forest <- function(ntrees)
{
  fit <- randomForest(as.factor(Spam) ~ ., data=spambase, subset = training, importance=TRUE,
                      ntree = ntrees)

  ypredict <- as.numeric(predict(fit, test))

  error <- mean((test$Spam - ypredict)^2)
}

error_rates_f <- sapply(number_of_trees, random_forest)
```



```

plot(error_rates_f,type = "b",main="Error Rate vs Tree Level", xlab= "Number of Trees", ylab= "Error",
      ylim=c(1.01,1.03), col="red", pch=19, cex=1)
## Question 2

mixture_model <- function(my_k)
{
  set.seed(1234567890)

  max_it <- 100 # max number of EM iterations
  min_change <- 0.1 # min change in log likelihood between two consecutive EM iterations
  N=1000 # number of training points
  D=10 # number of dimensions
  x <- matrix(nrow=N, ncol=D) # training data
  true_pi <- vector(length = 3) # true mixing coefficients
  true_mu <- matrix(nrow=3, ncol=D) # true conditional distributions
  true_pi=c(1/3, 1/3, 1/3)
  true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
  true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
  true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
  plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
  points(true_mu[2,], type="o", col="red")
  points(true_mu[3,], type="o", col="green")

  # Producing the training data

  for(n in 1:N) {

    k <- sample(1:3,1,prob=true_pi)
    for(d in 1:D) {
      x[n,d] <- rbinom(1,1,true_mu[k,d])
    }
  }

  K=my_k # number of guessed components
  z <- matrix(nrow=N, ncol=K) # fractional component assignments
  pi <- vector(length = K) # mixing coefficients
  mu <- matrix(nrow=K, ncol=D) # conditional distributions
  llik <- vector(length = max_it) # log likelihood of the EM iterations
  # Random initialization of the paramters
  pi <- runif(K,0.49,0.51)
  pi <- pi / sum(pi)
  for(j in 1:my_k) {
    mu[j,] <- runif(D,0.49,0.51)
  }
  pi
  mu

  for(it in 1:max_it)
  {
    if(K == 2)
    {
      plot(mu[1,], type="o", col="blue", ylim=c(0,1))
      points(mu[2,], type="o", col="red")
    }
  }

```

```

else if(K==3)
{
  plot(mu[1,], type="o", col="blue", ylim=c(0,1))
  points(mu[2,], type="o", col="red")
  points(mu[3,], type="o", col="green")
}
else
{
  plot(mu[1,], type="o", col="blue", ylim=c(0,1))
  points(mu[2,], type="o", col="red")
  points(mu[3,], type="o", col="green")
  points(mu[4,], type="o", col="yellow")
}

Sys.sleep(0.5)
# E-step: Computation of the fractional component assignment

# Bernoulli distribution
for (n in 1:N)
{
  prob_x=0

  for (k in 1:K)
  {
    prob_x=prob_x+prod( ((mu[k,]^x[n,])*((1-mu[k,])^(1-x[n,]))) ) * pi[k] #
  }

  for (k in 1:K)
  {
    z[n,k]=pi[k]*prod( ((mu[k,]^x[n,])*((1-mu[k,])^(1-x[n,]))) ) / prob_x
  }
}

#Log likelihood computation.

likelihood <-0
llik[it] <-0
for(n in 1:N)
{
  for (k in 1:K)
  {
    likelihood <-likelihood + z[n,k]*(pi[k]*prod( ((mu[k,]^x[n,])*((1-mu[k,])^(1-x[n,]))) ))
  }
  llik[it]<-llik[it]+log(likelihood)
}

cat("iteration: ", it, "log likelihood: ", llik[it], "\n")
flush.console()

```



```

# Stop if the log likelihood has not changed significantly
if (it > 1)
{
  if (llik[it]-llik[it-1] < min_change)
  {
    break
  }
}

#M-step: ML parameter estimation from the data and fractional component assignments

mu<- (t(z) %*% x) /colSums(z)

# N - Total no. of observations
pi <- colSums(z)/N
}

pi
mu
plot(llik[1:it], type="o")

mixture_model(2)
mixture_model(3)
mixture_model(4)
}

```