

# Basic concepts

## Introduction to Bayesian methods

### Lecture 1b

# Probabilities


- Frequentist tradition: probabilities derived from counts

**Example:** Tossing two dices

- $X$ =number on first dice
- $Y$ =number on second dice
- $p(x = X)$  frequency of observing  $X$
- $p(x = X, y = Y)$  frequency of observing  $X$  and  $Y$
- $p(x = X|y = Y)$  frequency of observing  $X$  given  $y=Y$



# Probability density



- $p(x \in [a, b]) = \int_a^b p(x)dx$
- $p(x) \geq 0, \int_{-\infty}^{+\infty} p(x)dx = 1$

# Probabilities

- **Laws of probabilities**

- Sum rule (compute **marginal** probability)

$$p(X) = \sum_Y p(X, Y)$$

$$p(X) = \int p(X, Y) dY$$

- Product rule

$$p(X, Y) = p(X|Y)p(Y)$$

Combination 1:

$$p(X) = \sum_Y p(X|Y)p(Y)$$

$$p(X) = \int p(X|Y)p(Y) dY$$



# Bayes theorem


- Combination 2:

## Bayes Theorem

$$p(Y|X) = \frac{p(X|Y)p(Y)}{p(X)}$$

$$p(Y|X) \propto p(X|Y)p(Y)$$

$$p(Y|X) = \frac{p(X|Y)p(Y)}{\int p(X|Y)p(Y)dY}$$



# Bayesian probabilities

- Probability reflects your knowledge (uncertainty) about a phenomenon → **subjective probabilities**
  - **Prior probability**  $p(w)$ , can be uninformative  $p(w) \propto 1$
  - Formulate a model, compute **likelihood**  $p(D|w)$
  - **Posterior probability**  $p(w|D)$ , after observing data
    - $p(w|D) \propto p(D|w)p(w)$
- Model parameters are considered as random variables
  - In real life, do not need to be random, but we model as random

# Basic ML ingredients

- Data  $D$ : observations

- Features  $X_1, \dots, X_p$
- Targets  $Y_1, \dots, Y_r$

Case	$X_1$	$X_2$	$Y$
1			
2			
...			

- Model  $P(x | w_1, \dots, w_k)$  or  $P(y | x, w_1, \dots, w_k)$

- Example: Linear regression  $p(y | x, w) = N(w_0 + w_1 x, \sigma^2)$

- Learning procedure (data  $\rightarrow$  get parameters  $\hat{w}$  or  $p(w | D)$ )

- Maximum likelihood, MAP, Bayes rule...

- Predict new data  $X^{new}$  by using the fitted model



# Probabilistic models

- A distribution  $p(x|w)$  or  $p(y|x, w)$
- Example:

- $x \sim \text{Bin}(n, \theta)$

$$p(x = k|n, \theta) = \binom{n}{k} \theta^k (1 - \theta)^{n-k}$$

- $y \sim N(\alpha_0 + \alpha_1 x, \sigma^2)$



Learn basic distributions and their properties → PRML, chapter 2!

Source: Wikipedia



# Fitting a model

- Given dataset  $D$  and model  $p(\mathbf{x}|\mathbf{w})$  or  $p(\mathbf{y}|\mathbf{x}, \mathbf{w})$ 
  - Frequentist approach: which combination of parameter values fits my data best?
  - Bayesian approach: parameters are random variables, all feasible values are acceptable
    - Different parameter values have different probabilities

# Fitting a model

- Frequentist principle: Maximum likelihood principle
  - Compute likelihood  $p(\mathbf{D} | w)$ 
$$p(\mathbf{D} | w) = \prod_{i=1}^n p(X_i | w) \text{ or } p(\mathbf{D} | w) = \prod_{i=1}^n p(Y_i | X_i, w)$$
  - Maximize the likelihood and find the optimal  $w^* \rightarrow$  they are the fitted values

## Remarks:


- Likelihood shows how much the chosen parameter value is proper for a specific model and the given data
- Normally **log-likelihood** is used in computations instead
- Other alternatives to ML exist...

# Fitting a model

**Example:** tossing a coin.

$$D = \{0,1,1,0,1,1,1,1,1,1,1\},$$

$$p(x = 1|\theta) = \theta, p(x = 0|\theta) = 1 - \theta$$



<http://cdn.toonvectors.com/images/35/10267/toonvectors-10267-940.jpg>

# Fitting a model

- Bayesian principle
  - Compute  $p(w|D)$  and then decide yourself what to do with this (for ex. MAP, mean, median)
- Use bayes theorem

$$p(w|D) = \frac{p(D|w)p(w)}{p(D)} \propto p(D|w)p(w)$$

- $p(D)$  is **marginal likelihood**
  - $p(D) = \int p(D|w)p(w)dw$  or
  - $p(D) = \sum_i p(D|w_i)p(w_i)$

**Example:** tossing a coin. Find  $p(\theta|D)$ , estimate various  $\theta^*$



# Fitting a model

- How to choose the prior?
  - Expert knowledge about the phenomenon
  - Forcing a model to have a certain structure
    - Example: decision trees: prior prefers smaller trees
  - Conjugacy [http://en.wikipedia.org/wiki/Conjugate\\_prior](http://en.wikipedia.org/wiki/Conjugate_prior)
    - Distribution of the posterior is the same type as the distribution of the likelihood or prior
- Prior is the most controversial about Bayesian methods, but
  - When  $N \rightarrow \infty$ , data overwhelms the prior

# Prediction

- **Plug-in estimation** (Frequentist and Bayesian)
  - Substitute the estimated  $w^*$  into  $p(\mathbf{x}|w)$  or  $p(y|\mathbf{x}^{new}, w)$
- **Bayesian model averaging**
  - Posterior predictive distribution:
    - $p(\mathbf{x}^{new}|D) = \int p(\mathbf{x}^{new}|w)p(w|D)dw$
    - $p(y|D) = \int p(y|w, \mathbf{x}^{new})p(w|D, \mathbf{x}^{new})dw$

# Black swan paradox

- In the coin example,  $p(x^{new} = 1) = \frac{k}{n}$  if MLE used
- If we made 3 attempts, no successes  $\rightarrow k=0$
- Does this mean  $p(x^{new} = 1) = 0$  ??
- Problem does not appear in Bayesian setting (posterior mean)

# Types of supervised models

- **Generative models:** model  $p(X|Y, w)$  and  $p(Y|w)$

- **Example:** k-NN classification

$$p(X = x|Y = C_i, K) = \frac{K_i}{N_i V}, p(C_i|K) = \frac{N_i}{N}$$

From Bayes Theorem,

$$p(Y = C_i|x, K) = \frac{K_i}{K}$$

- **Discriminative models:** model  $p(Y|X, w)$ ,  $X$  constant

- **Example:** logistic regression

- $p(Y = 1|\mathbf{w}, \mathbf{x}) = \frac{1}{1+e^{-\mathbf{w}^T \mathbf{x}}}$



# Generative vs Discriminative

- Generative can be used to generate new data
- Generative normally easier to fit (check Logistic vs K-NN)
- Generative: each class estimated separately → do not need to retrain when a new class added
- Discriminative models: can replace  $X$  with  $\phi(X)$  (preprocessing), method will still work
  - Not generative, distribution will change
- Generative: often make too strong assumptions about  $p(X|Y, w)$  → bad performance

# Bayesian decision theory

- Machine learning models estimate  $p(y|x)$  or  $p(y|x, \hat{w})$
- Transform probability into action  $\rightarrow$  which value to predict?  $\rightarrow$  decision step
  - $p(Y = Spam|x) = 0.83 \rightarrow$  do we move the mail to Junk?
  - What is more dangerous: deleting 1 non-spam mail or letting 1 spam mail enter Inbox?
- $\rightarrow$  **Loss function** or **Loss matrix**

# Loss matrix

- Costs of classifying  $Y = C_k$  to  $C_j$ :
  - Rows: true, columns: predicted

$$L = \|L_{ij}\|, i = 1, \dots, n, j = 1, \dots, n$$

- **Example 1:** 0/1-loss

$$L = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$$

- **Example 2:** Spam

$$L = \begin{pmatrix} 0 & 100 \\ 1 & 0 \end{pmatrix}$$

# Loss and decision


- Expected loss minimization

- $R_j$  : classify to  $C_j$

$$EL = \sum_k \sum_j \int_{R_j} L_{kj} p(\mathbf{x}, C_k) d\mathbf{x}$$

- Choose such  $R_j$  that  $EL$  is minimized
- Two classes

$$EL = \int_{R_1} L_{21} p(x, C_2) dx + \int_{R_2} L_{12} p(x, C_1) dx$$





# Loss and decision

- How to minimize  $EL$ ?
    - We are free to assign  $x$  to either  $R_1$  or  $R_2$
    - Assigning  $x$  to region with smallest  $L_{ij}p(x, C_i)$  will make  $EL$  smaller
  - $\rightarrow$  Rule:
    - $L_{21}p(x, C_1) > L_{12}p(x, C_2) \rightarrow$  predict  $y$  as  $C_1$
- $$\frac{p(C_1|x)}{p(C_2|x)} > \frac{L_{12}}{L_{21}} \rightarrow \text{predict } y \text{ as } C_1$$
- 0/1 Loss: **classify to the class which is more probable!**

# Loss and decision


- Continuous targets: squared loss

- Given a model  $p(x, y)$ ,  
minimize

$$EL = \int L(y, \hat{Y}(x)) p(x, y) dx dy$$

- Using **square loss**, the optimal is posterior mean

$$\hat{Y}(x) = \int y p(y|x) dy$$



# ROC curves

- Binary classification
- The choice of the threshold  $\hat{x} = \frac{L_{12}}{L_{21}}$  affects prediction  $\rightarrow$  what if we don't know the loss? Which classifier is better?
- **Confusion matrix**

	PREDICTED			
T R U E		1	0	Total
	1	TP	FN	$N_+$
	0	FP	TN	$N_-$

# ROC curves

- **True Positive Rates (TPR) = sensitivity = recall**

- Probability of detection of positives: TPR=1 positives are correctly detected

$$TPR = TP/N_+$$

- **False Positive Rates (FPR)**

- Probability of false alarm: system alarms (1) when nothing happens (true=0)

$$FPR = FP/N_-$$

- **Specificity**

$$Specificity = 1 - FPR$$

- **Precision**

$$Precision = \frac{TP}{TP + FP}$$



# ROC curves

- **ROC**=Receiver operating characteristics
- Use various thresholds, measure TPR and FPR
- Same FPR, higher TPR → better classifier
- Best classifier = greatest Area Under Curve (**AUC**)

