

Regression and regularization

Lecture 1d



Overview

- Linear regression
- Ridge Regression
- Lasso
- Variable selection

Simple linear regression

Model:

$$y \sim N(w_0 + w_1x, \sigma^2)$$

or

$$y = w_0 + w_1x + \epsilon, \\ \epsilon \sim N(0, \sigma^2)$$

or

$$p(y|x, w) = N(w_0 + w_1x, \sigma^2)$$


Terminology:

w_0 : intercept (or bias)

w_1 : regression coefficient

Response

The target responds directly and linearly to changes in the feature



Ordinary least squares regression (OLS)

Model:


$$y \sim N(\mathbf{w}^T \mathbf{x}, \sigma^2)$$

where

$$\mathbf{w} = \{w_0, \dots, w_d\}$$

$$\mathbf{x} = \{1, x_1, \dots, x_d\}$$

Why is "1" here?



The response variable responds directly and linearly to changes in each of the inputs

Ordinary least squares regression

Given data set D

Case	X_1	X_2			X_p	Y
1	x_{11}	x_{21}			x_{p1}	y_1
2	x_{12}	x_{22}			x_{p2}	y_2
3	x_{13}	x_{23}			x_{p3}	y_3
N	x_{1N}	x_{2N}			x_{pN}	y_N

Estimation: maximizing the likelihood

$$\hat{w} = \max_w p(D|w)$$

Is equivalent to minimizing

$$RSS(w) = \sum_{i=1}^n (Y_i - \mathbf{w}^T \mathbf{X}_i)^2$$

Matrix formulation of OLS regression

Optimality condition:

where $\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w}) = 0$

$$\mathbf{X} = \begin{pmatrix} 1 & x_{11} & x_{21} & x_{p1} \\ 1 & x_{12} & x_{22} & x_{p2} \\ \vdots & \vdots & \vdots & \vdots \\ 1 & x_{1N} & x_{2N} & x_{pN} \end{pmatrix} \quad \text{and} \quad \mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{pmatrix}$$

Parameter estimates and predictions

- Least squares estimates of the parameters

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- Predicted values

$$\hat{\mathbf{y}} = \mathbf{X} \hat{\mathbf{w}} = \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} = \mathbf{P} \mathbf{y}$$

- Linear regression belongs to the class of **linear smoothers**



Hat matrix

Why is it called so?

Degrees of freedom

Definition:

$$df(\hat{y}) = \frac{1}{\sigma^2} \sum_{i=1}^N \text{Cov}(\hat{y}_i, y_i)$$

- Larger **covariance** \rightarrow stronger connection \rightarrow model can approximate data better \rightarrow model more flexible (complex)
- For linear smoothers $\hat{Y} = S(X)Y$

$$df = \text{trace}(S)$$

- For linear regression, degrees of freedom is

$$df = \text{trace}(P) = p$$

Different types of features

- Interval variables
- Numerically coded ordinal variables
 - (small=1, medium=2, large=3)
- Dummy coded qualitative variables

Basis function expansion:

If $y = w_0 + w_1x_1 + w_2x_1^2 + w_3e^{-x_2} + \epsilon$,

Model becomes linear if to recompute:

$$\begin{aligned}\phi_1(x_1) &= x_1 \\ \phi_2(x_1) &= x_1^2 \\ \phi_3(x_1) &= e^{-x_2}\end{aligned}$$

Example of dummy coding:

$$x_1 = \begin{cases} 1, & \text{if Jan} \\ 0, & \text{otherwise} \end{cases}$$

$$x_2 = \begin{cases} 1, & \text{if Feb} \\ 0, & \text{otherwise} \end{cases}$$

.

.

.

$$x_{11} = \begin{cases} 1, & \text{if Nov} \\ 0, & \text{otherwise} \end{cases}$$

Basis function expansion

- In general $\phi_1(\dots)$ may be a function of several x components
- Having data given by \mathbf{X} , compute new data
- $$\Phi = \begin{pmatrix} 1 & \phi_1(x_{11}, \dots, x_{1p}) & \dots & \phi_p(x_{11}, \dots, x_{1p}) \\ \dots & \dots & \dots & \dots \\ 1 & \phi_1(x_{n1}, \dots, x_{np}) & \dots & \phi_p(x_{n1}, \dots, x_{np}) \end{pmatrix}$$
- If doing a basis function in a model, replace \mathbf{X} by Φ everywhere where \mathbf{X} is used:

$$\hat{\mathbf{y}} = \Phi(\Phi^T \Phi)^{-1} \Phi^T \mathbf{y}$$

Linear regression in R

- `fit=lm(formula, data, subset, weights,...)`
 - **data** is the data frame containing the predictors and response values
 - **formula** is expression for the model
 - **subset** which observations to use (training data)?
 - **weights** should weights be used?

fit is object of class **lm** containing various regression results.

- Useful functions (many are generic, used in many other models)
 - Get details about the particular function by `"."`, for ex. `predict.lm`

`summary(fit)`

`predict(fit, newdata, se.fit, interval)`

`coefficients(fit)` # model coefficients

`confint(fit, level=0.95)` # CIs for model parameters

`fitted(fit)` # predicted values

`residuals(fit)` # residuals

An example of ordinary least squares regression

```
mydata=read.csv2("Bilexempel.csv")
fit1=lm(Price~Year, data=mydata)
summary(fit1)
fit2=lm(Price~Year+Mileage+Equipment,
data=mydata)
summary(fit2)
```

Response variable:

Requested price of used Porsche cars
(1000 SEK)

```
> summary(fit1)
```

Call:

```
lm(formula = Price ~ Year, data = mydata)
```

Residuals:

Min	1Q	Median	3Q	Max
-167683	-14683	20056	35933	72317

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-78161027	8448038	-9.252	6.00e-13 ***
Year	39246	4226	9.288	5.25e-13 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 57270 on 57 degrees of freedom
Multiple R-squared: 0.9921, Adjusted R-squared: 0.5952
F-statistic: 86.26 on 1 and 57 DF, p-value: 5.248e-13

Inputs:

X_1 = Manufacturing year

X_2 = Milage (km)

X_4 = Equipment (0 or 1)

An example of ordinary least squares regression

```
> summary(fit2)
```

Call:

```
lm(formula = Price ~ Year + Mileage + Equipment, data = mydata)
```

Residuals:

Min	1Q	Median	3Q	Max
-66223	-10525	-739	14128	65332

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-2.083e+07	6.309e+06	-3.302	0.00169	**
Year	1.062e+04	3.154e+03	3.366	0.00139	**
Mileage	-2.077e+00	2.022e-01	-10.269	2.14e-14	***
Equipment	5.790e+04	1.041e+04	5.563	8.08e-07	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 29270 on 55 degrees of freedom

Multiple R-squared: 0.8997, Adjusted R-squared: 0.8942

F-statistic: 164.5 on 3 and 55 DF, p-value: < 2.2e-16


An example of ordinary least squares regression

- Prediction

```
fitted <- predict(fit1, interval =  
"confidence")
```

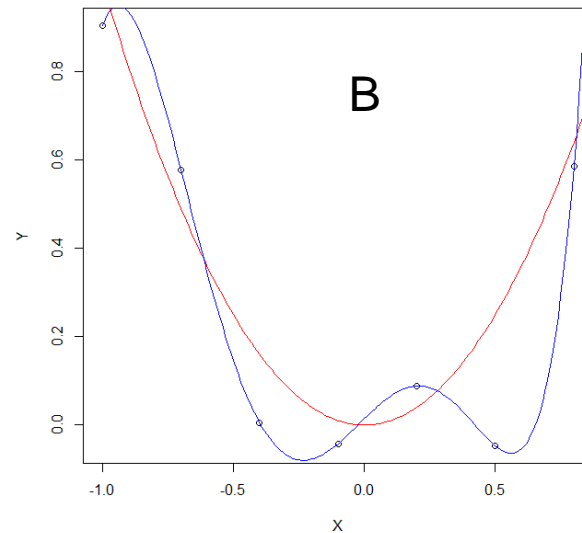

```
# plot the data and the fitted line  
attach(mydata)  
plot(Year, Price)  
lines(Year, fitted[, "fit"])
```

```
# plot the confidence bands  
lines(Year, fitted[, "lwr"], lty = "dotted",  
col="blue")  
lines(Year, fitted[, "upr"], lty = "dotted",  
col="blue")  
detach(mydata)
```



Ridge regression

- Problem: linear regression can overfit:
 - Take $Y := Y, X_1 = X, X_2 = X^2, \dots, X_p = X^p \rightarrow$ polynomial model, fit by linear regression
 - High degree of polynomial leads to overfitting.



Ridge regression

- **Idea:** Keep all predictors but shrink coefficients to make model less complex

minimize $-\log\text{likelihood} + \lambda_0 \|w\|_2^2$

→ **l_2 regularization**

- Given that model is Gaussian, we get **Ridge regression:**

$$\hat{w}^{\text{ridge}} = \operatorname{argmin} \left\{ \sum_{i=1}^N (y_i - w_0 - w_1 x_{1j} - \dots - w_p x_{pj})^2 + \lambda \sum_{j=1}^p w_j^2 \right\}$$

- $\lambda > 0$ is **penalty factor**

Ridge regression

Equivalent form

$$\hat{w}^{ridge} = \operatorname{argmin} \sum_{i=1}^N (y_i - w_0 - w_1 x_{1j} - \dots - w_p x_{pj})^2$$

subject to $\sum_{j=1}^p w_j^2 \leq s$

Solution

$$\hat{w}^{ridge} = (X^T X + \lambda I)^{-1} X^T y$$

$$\hat{y} = X\hat{w} = X(X^T X + \lambda I)^{-1} X^T y = Py$$

Hat matrix

How do we
compute degrees
of freedom here?

Ridge regression

Properties

- Extreme cases:
 - $\lambda = 0$ usual linear regression (no shrinkage)
 - $\lambda = +\infty$ fitting a constant ($w = 0$ except of w_0)
- When input variables are orthogonal (not realistic), $X^T X = I \rightarrow$
$$\hat{\mathbf{w}}^{\text{ridge}} = \frac{1}{1+\lambda} \mathbf{w}^{\text{linreg}} \rightarrow \text{coefficients are equally shrunk}$$
- **Ridge regression is particularly useful if the explanatory variables are strongly correlated to each other.**
 - Correlated variables often correspond large $w \rightarrow$ shrunk
- Degrees of freedom decrease when λ increases
 - $\lambda = 0 \rightarrow d.f. = p$

Ridge regression

Properties

- Shrinking enables estimation of regression coefficients even if the number of parameters exceeds the number of cases!
($X^T X + \lambda I$ is always nonsingular)
 - Compare with linear regression
- How to estimate λ ?
 - cross-validation

Ridge regression

- Bayesian view

- Ridge regression is just a special form of Bayesian Linear Regression with constant σ^2 :

$$\begin{aligned} \mathbf{y} &\sim N(\mathbf{y} | \mathbf{w}_o + \mathbf{X}\mathbf{w}, \sigma^2 \mathbf{I}) \\ \mathbf{w} &\sim N\left(\mathbf{0}, \frac{\sigma^2}{\lambda} \mathbf{I}\right) \end{aligned}$$

Theorem MAP estimate to the Bayesian Ridge is equal to solution in frequentist Ridge

$$\hat{\mathbf{w}}^{ridge} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

- In Bayesian version, we can also make inference about λ

Ridge regression

Example Computer Hardware Data Set : performance measured for various processors and also

- Cycle time
- Memory
- Channels
- ...

Build model predicting performance




Ridge regression

- R code: use package **glmnet** with $\alpha=0$ (Ridge regression)
- Seeing how Ridge converges

```
data=read.csv("machine.csv", header=F)
covariates=scale(data[,3:8])
response=scale(data[, 9])
```

```
model0=glmnet(as.matrix(covariates),
response, alpha=0,family="gaussian")
plot(model0, xvar="lambda", label=TRUE)
```




Ridge regression

- Choosing the best model by cross-validation:

```
model=cv.glmnet(as.matrix(covariates),  
response, alpha=0,family="gaussian")  
model$lambda.min  
plot(model)  
coef(model, s="lambda.min")
```

```
> coef(model, s="lambda.min")  
7 x 1 sparse Matrix of class "dgCMatrix"  
1  
(Intercept) -4.530442e-17  
v3          3.420739e-02  
v4          3.085696e-01  
v5          3.403839e-01  
v6          1.593470e-01  
v7          5.489116e-02  
v8          1.970982e-01
```



```
> model$lambda.min  
[1] 0.046
```

Ridge regression

- How good is this model in prediction?

```
ind=sample(209, floor(209*0.5))
data1=scale(data[,3:9])
train=data1[ind,]
test=data1[-ind,]

covariates=train[,1:6]
response=train[, 7]
model=cv.glmnet(as.matrix(covariates), response, alpha=1,family="gaussian",
lambda=seq(0,1,0.001))
y=test[,7]
ynew=predict(model, newx=as.matrix(test[, 1:6]), type="response")
```

#Coefficient of determination

```
sum((ynew-mean(y))^2)/sum((y-mean(y))^2)
```

```
sum((ynew-y)^2)
```

Note that data are so small so numbers
change much for other train/test

```
> sum((ynew-mean(y))^2)/sum((y-mean(y))^2)
[1] 0.5438148
> sum((ynew-y)^2)
[1] 18.04988
> |
```


LASSO

- **Idea:** Similar idea to Ridge
- Minimize minus loglikelihood plus **linear** penalty factor $\rightarrow I_1$ **regularization**
 - Given that model is Gaussian, we get **LASSO** (least absolute shrinkage and selection operator):

$$\hat{w}^{lasso} = \operatorname{argmin} \left\{ \sum_{i=1}^N (y_i - w_0 - w_1 x_{1j} - \dots - w_p x_{pj})^2 + \lambda \sum_{j=1}^p |w_j| \right\}$$

- $\lambda > 0$ is **penalty factor**



LASSO

- Equivalently

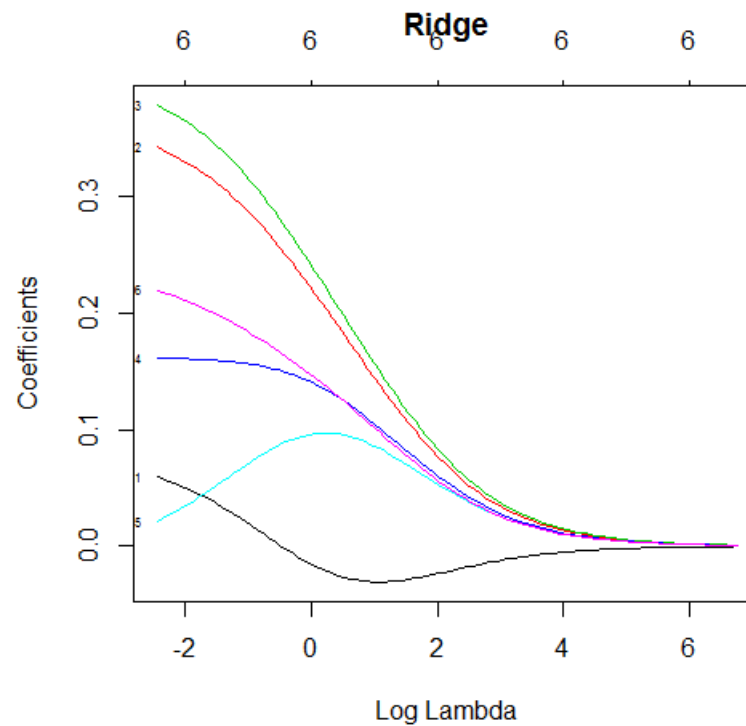

$$\hat{w}^{lasso} = \operatorname{argmin} \sum_{i=1}^N (y_i - w_0 - w_1 x_{1j} - \dots - w_p x_{pj})^2$$

subject to $\sum_{j=1}^p |w_j| \leq s$

LASSO vs Ridge

- **LASSO yields sparse solutions!**

Example Computer hardware data



LASSO vs Ridge

- Only 5 variables selected by LASSO

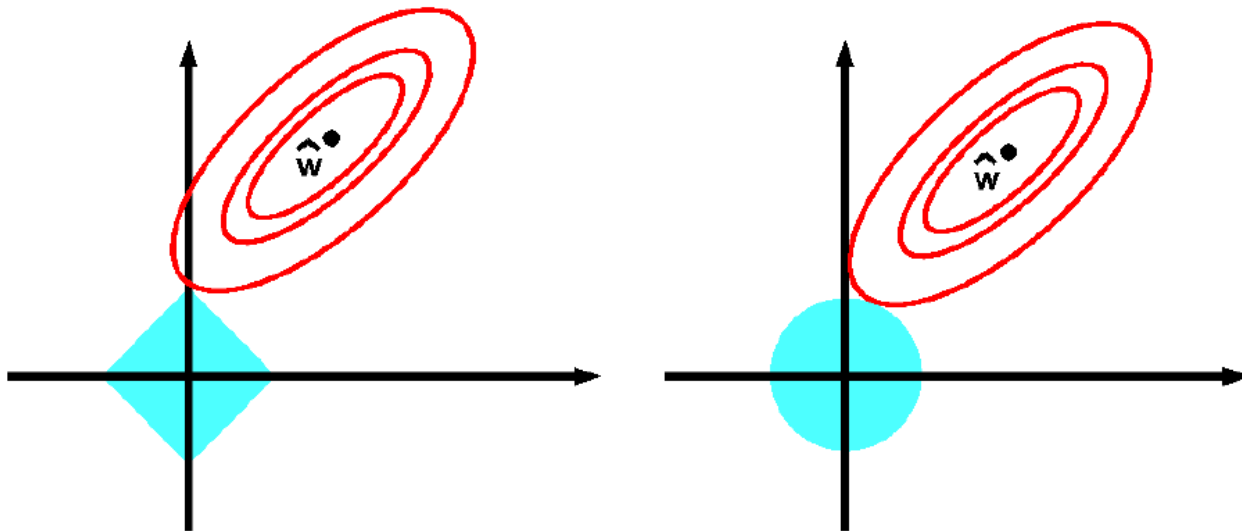
```
> coef(model, s="lambda.min")  
7 x 1 sparse Matrix of class "dgCMatrix"
```

```
      1  
(Intercept) -5.091825e-17  
v3           6.350488e-02  
v4           3.578607e-01  
v5           4.033670e-01  
v6           1.541329e-01  
v7           .  
v8           2.287134e-01  
> l
```

```
> sum((ynew-mean(y))^2)/sum((y-mean(y))^2)  
[1] 0.5826904  
> sum((ynew-y)^2)  
[1] 16.63756
```


LASSO vs Ridge

- Why Lasso leads to sparse solutions?
 - Feasible area for Ridge is a circle (2D)
 - Feasible area for LASSO is a polygon (2D)



LASSO properties

- **Lasso is widely used when $p \gg n$**
 - Linear regression breaks down when $p > n$
 - Application: DNA sequence analysis, Text Prediction

- When inputs are orthonormal,

$$\hat{w}_i^{\text{lasso}} = \text{sign}(w_i^{\text{linreg}}) \left(|w_i^{\text{linreg}}| - \frac{\lambda}{2} \right)_+$$

- No explicit formula for \hat{w}^{lasso}
 - Optimization algorithms used

**Coding in R: use
glmnet() with
alpha=1**

Variable selection

- .. Or “Feature selection”

Often, we do not need all features available in the data to be in the model

Reasons:



- Model can become overfitted (recall polynomial regression)
- Large number of predictors → model is difficult to use and interpret

Variable selection

Alternative 1: Variable subset selection

- Best subset selection:
 - Consider different subsets of the full set of features, fit models and evaluate their quality
 - Problem: computationally difficult for p around 30 or more
 - How to choose the best model size? Some measure of predictive performance normally used (ex. AIC).
- Forward and Backward stepwise selection
 - Starts with 0 features (or full set) and then adds a feature (removes feature) that most improves the measure selected.
 - Can handle large p quickly
 - Does not examine all possible subsets (not the “best”)

RSS and MSE depend on k



Variable selection in R

- Use stepAIC() in MASS

```
library(MASS)
fit <- lm(V9~.,data=data.frame(data1))
step <- stepAIC(fit, direction="both")
step$anova
summary(step)
```

```
Call:
lm(formula = V9 ~ V3 + V4 + V5 + V6 + V8, data = data.frame(data1))
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-1.20232 -0.15512  0.03579  0.16567  2.42280
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -5.783e-17  2.574e-02   0.000   1.0000
V3           7.948e-02  2.826e-02   2.813   0.0054 **
V4           3.661e-01  4.312e-02   8.490  4.34e-15 ***
V5           4.055e-01  4.664e-02   8.695  1.18e-15 ***
V6           1.591e-01  3.394e-02   4.687  5.07e-06 ***
V8           2.360e-01  3.356e-02   7.031  3.06e-11 ***
```

```
> step <- stepAIC(fit, direction="both")
```

```
Start:  AIC=-405.35
```

```
V9 ~ V3 + V4 + V5 + V6 + V7 + V8
```

	Df	Sum of Sq	RSS	AIC
- V7	1	0.0139	28.117	-407.25
<none>			28.103	-405.35
- V3	1	1.0819	29.185	-399.46
- V6	1	2.9385	31.041	-386.57
- V8	1	6.3150	34.418	-364.99
- V4	1	9.7492	37.852	-345.11
- V5	1	10.4837	38.586	-341.09

```
Step:  AIC=-407.25
```

```
V9 ~ V3 + V4 + V5 + V6 + V8
```

	Df	Sum of Sq	RSS	AIC
<none>			28.117	-407.25
+ V7	1	0.0139	28.103	-405.35
- V3	1	1.0958	29.212	-401.26
- V6	1	3.0431	31.160	-387.77
- V8	1	6.8472	34.964	-363.70
- V4	1	9.9840	38.101	-345.74
- V5	1	10.4713	38.588	-343.08