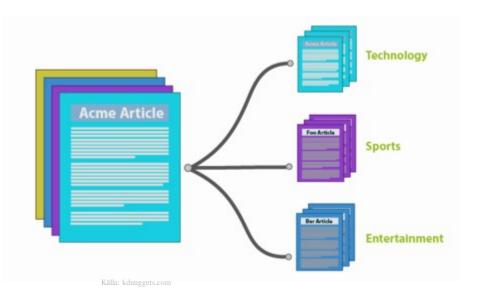


Naive Bayes classifiers: motivation

- Consider n labeled text documents
 - Y = {0,1}, 0 = "Science fiction", 1 = "Comedy"
 - $-X = \{X_1, ... X_{100}\}$ does the document contain the keyword (0=No,1=Yes)
 - X_1 corr. "space", X_2 corr. "fun",...



Want to classify a new document

Naive Bayes classifiers: motivation

Chance of observing a given combination of words in science fiction

Idea: use Bayes classifier

Payes classifier
$$p(Y = y|X) = \frac{P(X|Y = y)P(Y = y)}{\sum_{j} P(X|Y = y_{j})P(Y = y_{j})}$$

Proportion of science fiction documents

Naive Bayes classifiers: motivation

Attempt 1:

- Model $P(X = (x_1, ... x_p) | Y = y_i)$ and $P(Y = y_i)$ as unknown parameters
- Use data to derive those with Maximum Likelihood
- Classify by use of the posterior distribution
- How many parameters?
 - How many different combinations of X? 2^p
 - Amount of $P(X = (x_1, ... x_p) | Y = y_i)$ is $2 * 2^p 2$
 - Probabilities for each Y sum up to one
- If p = 100, 10^{30} parameters need to be estimated \rightarrow ouch!

Naive Bayes classifiers

Naive Bayes assumption: conditional independence

$$P(X = (x_1, ... x_p) | Y = y) = \prod_{i=1}^p P(X_i = x_i | Y = y)$$

How many parameters now?

$$-P(X_i = x_i | Y = y), i = 1, ... p, x_i = \{0,1\}, y = \{0,1\} \ 2 * p$$

- Is Naive Bayes assumption always valid?
 - P(Space,ship|SciFi)= P(Space|SciFi)*P(Ship|SciFi) ?

Naive Bayes classifiers - discrete inputs

- Given $D = \{(X_{m1}, ... X_{mp}, Y_m), m = 1, ... n\}$
- Assume $X_i \in \{x_1, ... x_I\}, i = 1, ... p, Y \in \{y_1 ... y_K\}$
- Denote $\theta_{ijk} = p(X_i = x_j | Y = y_k)$
 - How many parameters? (J-1)Kp
- Denote $\pi_k = p(Y = y_k)$
- Maximum likelihood: assume θ_{ijk} and π_k are constants

$$- \hat{\theta}_{ijk} = \frac{\#\{X_i = x_j \& Y = y_k\}}{\#\{Y = y_k\}}$$

$$- \hat{\pi}_k = \frac{\#\{Y = y_k\}}{n}$$

- Classification using 0-1 loss: $\hat{Y} = \arg \max_{y} p(Y = y|X)$

Naive Bayes classifiers - discrete inputs

- Example Loan decision
 - Classify a person: Home Owner=No, Single=Yes

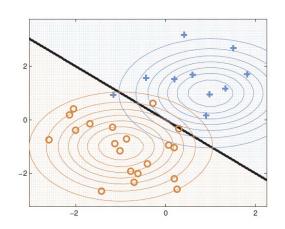
	10.00	7.		-	
Tid	Home Owner	Marital Status	Annual Income	Defaulted Borrower No No No No Yes	
1	Yes	Single	125K		
2	No	Married	100K		
3	No	Single	70K		
4	Yes	Married	120K		
5	No	Divorced	95K		
6	No	Married	60K	No	
7	Yes	Divorced	220K	No	
8	No	Single	85K	Yes	
9	No	Married	75K	No	
10	No	Single	90K	Yes	

Naive Bayes – continuous inputs

- X_i are continuous
- Assumption A: $x_i | y = C$ are univariate Gaussian

$$- p(x_j|y = C_i, \theta) = N(x_j|\mu_{ij}, \sigma_{ij}^2)$$

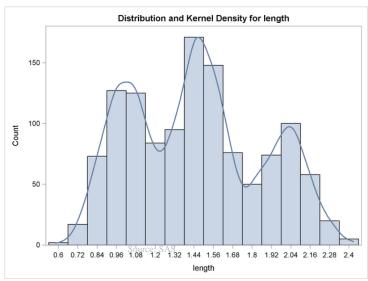
- Therefore $p(x|y = C_i, \theta) = N(x|\mu_i, \Sigma_i)$
 - Σ_i =diag($\sigma_{i1}^2, ..., \sigma_{ip}^2$)



- Naive bayes is a special case of LDA (given A)
 - → MLE are means and variances (per class)

Naive Bayes – continuous inputs

- Assumption B: $p(x_j|y=C)$ are unknown functions of x_j that can be estimated from data
 - Nonparametric density estimation (kernel for ex.)
 - 1. Estimate $p(X_i = x_i | Y = y_k)$ using nonparametric methods
- 2. Estimate $p(Y = y_k)$ as class proportions
- 3. Use Bayes rule and 0-1 loss to classify



Naive Bayes in R

naiveBayes in package e1071

Example: Satisfaction of householders with their present housing circumstances

```
library(MASS)
library(e1071)
n=dim(housing)[1]
ind=rep(1:n, housing[,5])
housing1=housing[ind,-5]

fit=naiveBayes(Sat~., data=housing1)
fit

Yfit=predict(fit, newdata=housing1)
table(Yfit,housing1$Sat)
```

```
> table(Yfit,housing1$Sat)
```

```
Yfit Low Medium High
Low 294 162 144
Medium 20 23 20
High 253 261 504
```

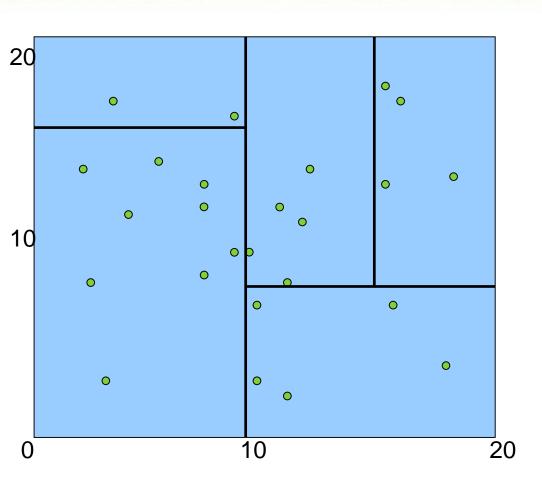
Decision trees

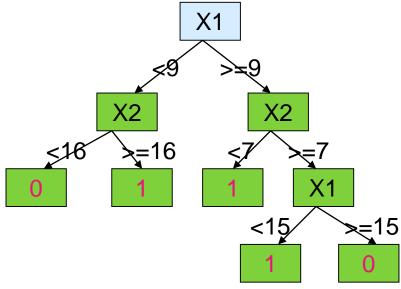
Idea

Split the domain of predictor set into the set of hypercubes (rectangles, cubes) and define the target value to be constant within each hypercube

- Regression trees:
 - Target is a continuous variable
- Classification trees
 - Target is a class (qualitative) variable

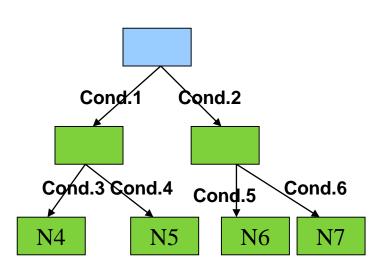
Classification tree toy example



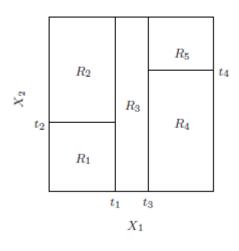


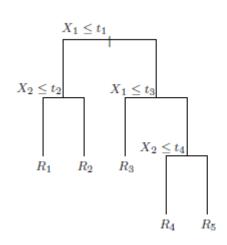
Definitions

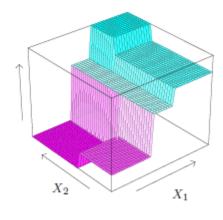
- Root node
- Nodes
- Leaves (terminal nodes)
- Parent node, child node
- Decision rules
- A value is assigned to the leaves



Regression tree toy example







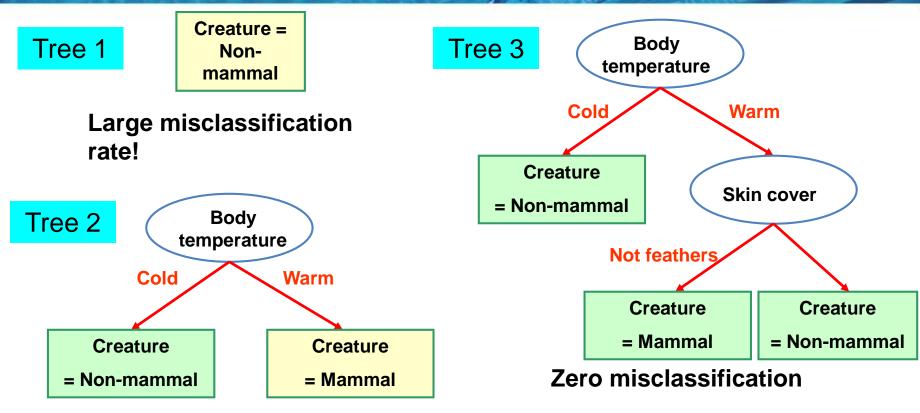
A classification problem

Create a classification tree that would describe the following patterns

ID x1 x2 x3 x4 x5 x6 x7 y

	Body			Aquatic	Aerial			
Name	temperature	Skin cover	Gives birth	creature	creature	Has legs	Hibernates	Class label
human	warm-blooded	hair	yes	no	no	yes	no	mammal
python	cold-blooded	scales	no	no	no	no	yes	non-mammal
salmon	cold-blooded	scales	no	yes	no	no	no	non-mammal
whale	warm-blooded	hair	yes	yes	no	no	no	mammal
frog	cold-blooded	none	no	semi	no	yes	yes	non-mammal
komodo	cold-blooded	scales	no	no	no	yes	no	non-mammal
bat	warm-blooded	hair	yes	no	yes	yes	yes	mammal
pigeon	warm-blooded	feathers	no	no	yes	yes	no	non-mammal
cat	warm-blooded	fur	yes	no	no	yes	no	mammal
shark	cold-blooded	scales	yes	yes	no	no	no	non-mammal
turtle	cold-blooded	scales	no	semi	no	yes	no	non-mammal
penguin	warm-blooded	feathers	no	semi	no	yes	no	non-mammal
porcupine	warm-blooded	quills	yes	no	no	yes	yes	mammal
eel	cold-blooded	scales	no	yes	no	no	no	non-mammal
salamander	cold-blooded	none	no	semi	no	yes	yes	non-mammal

Several solutions



A lower misclassification rate

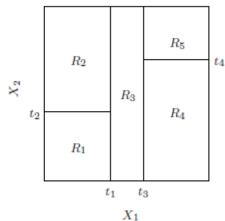
Green boxes represent pure nodes =nodes where observed values are the same

Decision trees

- A tree $T = \langle r_i, s_{r_i}, R_j, i = 1 ... S, j = 1 ... L \rangle$
 - $-x_{r_i} \leq s_{r_i}$ splitting rules (conditions), S- their amount
 - $-R_i$ -terminal nodes, L- their amount
 - labels μ_i in each terminal node

Model:

- Y|T for R_j comes from exponential family with mean μ_j
- Fitting by MLE:
 - Step 1: Finding optimal tree
 - Step 2: Finding optimal labels in terminal nodes



Decision trees

Example:

- Normal model leads to regression trees
 - Objective: MSE
- Multinoulli model leads to classification trees
 - Objective: cross-entropy (deviance)

Classification trees

- Target is categorical
- Classification probability $p_{mk} = p(Y = k | X \in R_m)$ is estimated for every class in a node
- How to estimate p_{mk} for class k and node R_m ?

Class proportions

$$\hat{p}_{mk} = \frac{1}{N_m} \sum_{x_i \in R_m} I(y_i = k)$$

• For any node (leave), a label can be assigned

$$k(m) = \arg \max_{k} \hat{p}_{mk}$$

Classification trees

- Impurity measure $Q(R_m)$
 - $-R_m$ is a tree node (region)
 - Node can be split unless it is pure

Misclassification error:
$$\frac{1}{N_m} \sum_{i \in R_m} I(y_i \neq k(m)) = 1 - \hat{p}_{mk(m)}$$

Gini index:
$$\sum_{k \neq k'} \hat{p}_{mk} \hat{p}_{mk'} = \sum_{k=1}^{K} \hat{p}_{mk} (1 - \hat{p}_{mk})$$

Cross-entropy or deviance:
$$-\sum_{k=1}^{K} \hat{p}_{mk} \log \hat{p}_{mk}$$
.

• Note: In many sources, deviance is $Q(R_m) N(R_m)$

Example: Cross –entropy is MLE of $Y_j | T \sim Multinomial(p_{j1}, ... p_{jc})$

Fitting regression trees: CART

Step 1: Finding optimal tree: grow the tree in order to minimize global objective

- 1. Let C_0 be a hypercube containing all observations
- 2. Let queue $C=\{C_0\}$
- 3. Pick up some C_i from C and find a variable X_j and value s that split C_j into two hypercubes

and solve

$$R_1(j,s) = \{X|X_j \le s\}$$
 and $R_2(j,s) = \{X|X_j > s\}$
$$\min_{j,s} [N_1Q(R_1) + N_2Q(R_2)]$$

- 4. Remove C_i from C and add R₁ and R₂
- Repeat 3-4 as many times as needed (or until each cube has only 1 observation)

CART: comments

Greedy algorithm (optimal tree is not found)

- The largest tree will interpolate the data → large trees = overfitting the data
- Too small trees=underfitting (important structure may not be captured)
- Optimal tree length?

Optimal trees

Postpruning

Weakest link pruning:

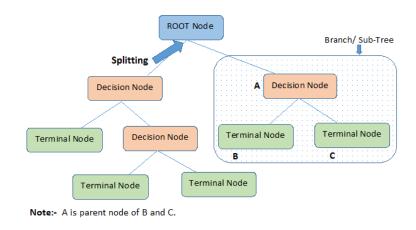
- 1. Merge two leaves that have smallest N(parent)*Q(parent)-N(leave1)Q(leave1)-N(leave2)Q(leave2)
- 2. For the current tree T, compute

$$I(T) = \sum_{R_i \in leaves} N(R_i)Q(R_i) + \alpha |T|$$

|T| =#leaves

- 3. Repeat 1-2 until the tree with one leave is obtained
- 4. Select the tree with smallest I(T)

How to find the optimal α ? Cross validation!



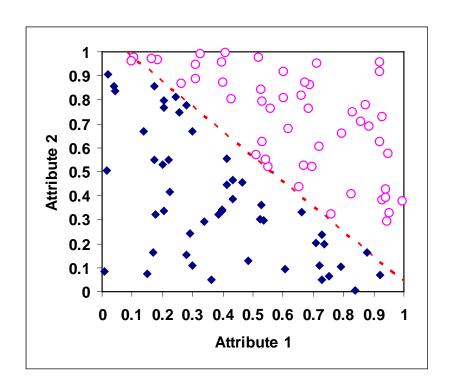
Decision trees: comments

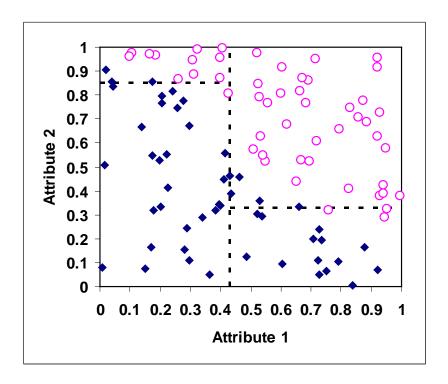
- Similar algorithms work for regression trees replace $N\cdot Q(R)$ by SSE(R)
- Easy to interpret
- Easy to handle all types of predictors in one model
- Automatic variable selection
- Relatively robust to outliers
- Handle large datasets
- Trees have high variance: a small change in response

 totally different tree
- Greedy algorithms → fit may be not so good
- Lack of smoothness

Decision trees: issues

Large trees may be needed to model an easy system:





- **tree** package
 - Alternative: rpart

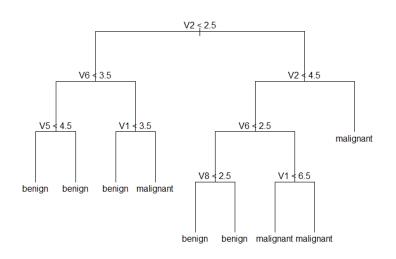
```
tree(formula, data, weights, control, split = c("deviance", "gini"), ...)
print(), summary(), plot(), text()
```

Example: breast cancer as a function av biological measurements

```
library(tree)
n=dim(biopsy)[1]
fit=tree(class~., data=biopsy)
plot(fit)
text(fit, pretty=0)
fit
summary(fit)
```

Adjust the splitting in the tree with control parameter (leaf size for ex)

```
> fit
node), split, n, deviance, yval, (yprob)
      * denotes terminal node
 1) root 683 884.400 benign ( 0.650073 0.349927 )
   2) V2 < 2.5 418 108.900 benign ( 0.971292 0.028708 )
     4) V6 < 3.5 395 25.130 benign ( 0.994937 0.005063 )
       8) V5 < 4.5 389 0.000 benign ( 1.000000 0.000000 ) *
       9) V5 > 4.5 6 7.638 benign ( 0.666667 0.333333 ) *
     5) V6 > 3.5 23 31.490 benign ( 0.565217 0.434783 )
     10) V1 < 3.5 11 0.000 benign ( 1.000000 0.000000 ) *
      11) V1 > 3.5 12 10.810 malignant ( 0.166667 0.833333 ) *
   3) V2 > 2.5 265 217.900 malignant ( 0.143396 0.856604 )
     6) V2 < 4.5 90 120.300 malignant ( 0.388889 0.611111 )
     12) V6 < 2.5 30 27.030 benign ( 0.833333 0.166667 )
        24) V8 < 2.5 19 0.000 benign ( 1.000000 0.000000 ) *
        25) V8 > 2.5 11 15.160 benign ( 0.545455 0.454545 ) *
      13) V6 > 2.5 60 54.070 malignant ( 0.166667 0.833333 )
        26) V1 < 6.5 28 35.160 malignant ( 0.321429 0.678571 ) *
        27) V1 > 6.5 32 8.900 malignant ( 0.031250 0.968750 ) *
     7) V2 > 4.5 175 30.350 malignant ( 0.017143 0.982857 ) *
```



> summary(fit)

```
Classification tree:
tree(formula = class ~ ., data = biopsy)
Variables actually used in tree construction:
[1] "V2" "V6" "V5" "V1" "V8"
Number of terminal nodes: 9
Residual mean deviance: 0.1603 = 108 / 674
Misclassification error rate: 0.03221 = 22 / 683
```

Misclassification results

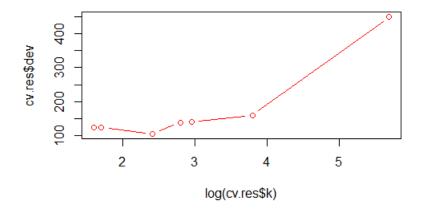
```
Yfit=predict(fit, newdata=biopsy, type="class")
table(biopsy$class,Yfit)
```

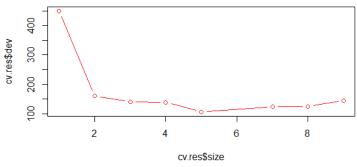
- Selecting optimal tree by penalizing
 - Cv.tree()

```
set.seed(12345)
ind=sample(1:n, floor(0.5*n))
train=biopsy[ind,]
valid=biopsy[-ind,]

fit=tree(class~., data=train)
set.seed(12345)
cv.res=cv.tree(fit)
plot(cv.res$size, cv.res$dev, type="b", col="red")
plot(log(cv.res$k), cv.res$dev,
type="b", col="red")
```

What is optimal number of leaves?





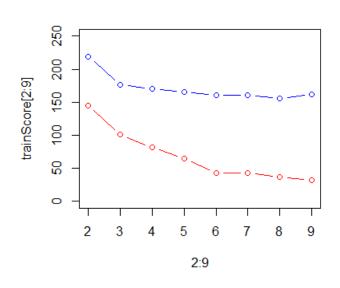
Selecting optimal tree by train/validation

```
fit=tree(class~., data=train)

trainScore=rep(0,9)

testScore=rep(0,9)

for(i in 2:9) {
    prunedTree=prune.tree(fit,best=i)
    pred=predict(prunedTree, newdata=valid,
    type="tree")
        trainScore[i]=deviance(prunedTree)
        testScore[i]=deviance(pred)
}
plot(2:9, trainScore[2:9], type="b", col="red",
    ylim=c(0,250))
points(2:9, testScore[2:9], type="b", col="blue")
```



What is optimal number of leaves?

• Final tree: 5 leaves

```
finalTree=prune.tree(fit, best=5)
Yfit=predict(finalTree, newdata=valid,
type="class")
table(valid$class,Yfit)
```