

Principal component regression. Uncertainty estimation

Lecture 2d

Principle component regression

Step 1: Compute principal components $v_1 \dots v_M$

Step 2: Compute derived data as $z_i = Xv_i$

Step 3: Compute linear regression using data set Z with columns $z_1 \dots z_M$, Y

The result

$$\hat{Y}_{pcr} = \sum_{m=1}^M \frac{\langle z_m, Y \rangle}{\langle z_m, z_m \rangle} z_m$$

Coefficients

$$\hat{\beta}_{pcr} = \sum_{m=1}^M \frac{\langle z_m, Y \rangle}{\langle z_m, z_m \rangle} v_m$$

PCR: comments

- Result depends on the scaling of features → standardize the original data
- If $M < p \rightarrow$ reduced regression
- If $M = p \rightarrow$ The fitted response will be the same
- Ridge regression shrinks the coefficients along the principal components, principal components regression discards $p - M$ smallest components

Partial Least Squares Regression (PLS)

- Idea with PLS is to find $M < p$ orthogonal directions (as in PCR). The difference:

PCR

$$\begin{aligned} & \max_{\alpha} \text{Var}(\mathbf{X}\alpha) \\ & \text{subject to } \|\alpha\| = 1, \alpha^T \mathbf{S} \mathbf{v}_{\ell} = 0, \ell = 1, \dots, m-1, \end{aligned}$$

PLS

$$\begin{aligned} & \max_{\alpha} \text{Corr}^2(\mathbf{y}, \mathbf{X}\alpha) \text{Var}(\mathbf{X}\alpha) \\ & \text{subject to } \|\alpha\| = 1, \alpha^T \mathbf{S} \hat{\mathbf{v}}_{\ell} = 0, \ell = 1, \dots, m-1. \end{aligned}$$

Partial least squares regression (PLS)

Step 1: Standardize features to mean zero and variance one

Step 2: Compute the first derived feature by setting

$$\mathbf{z}_1 = \sum_{j=1}^p \varphi_{1j} \mathbf{x}_j$$

where the φ_{1j} is projection of \mathbf{Y} on \mathbf{x}_j

Step 3: Orthogonalize $\mathbf{x}_1 \dots \mathbf{x}_m$ with respect to \mathbf{z}_1

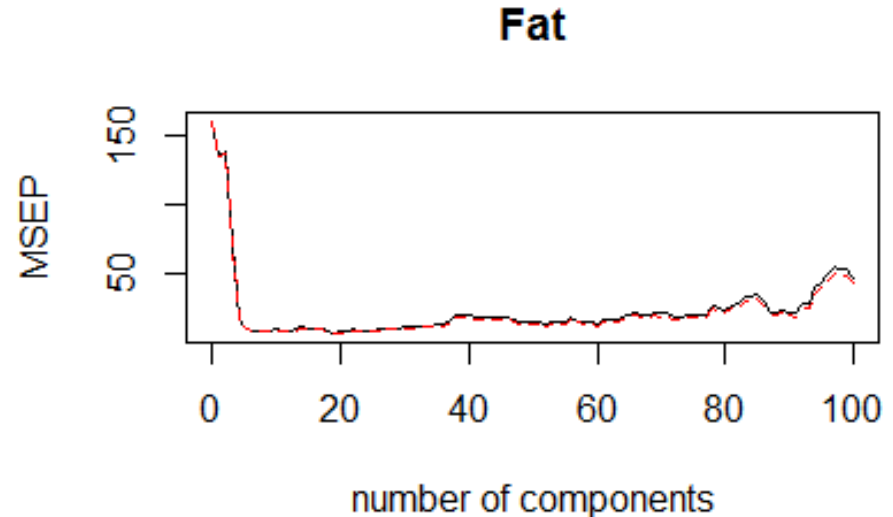
Step 4: repeat from step 2 and find $\mathbf{z}_2 \dots \mathbf{z}_M$

Step 5: Compute regression of \mathbf{Y} on $\mathbf{z}_1 \dots \mathbf{z}_M$

PCR and PLS: R

- Package **pls**
- **PCR**: `pcr(formula, ncomp, data, scale = FALSE, validation = c("none", "CV", "LOO"...))`
- **PLS**: `plsr(...)`

```
predictors=paste("Channel", 1:100,  
sep="")  
f=formula(paste("Fat ~ ",  
paste(predictors, collapse=" + ")))  
  
set.seed(12345)  
pcr.fit=pcr(f, data=train,  
validation="CV")  
summary(pcr.fit)  
validationplot(pcr.fit, val.type="MSEP")
```



PCR and PLS: R

```
> summary(pcr.fit)
```

```
Data: X dimension: 150 100
```

```
Y dimension: 150 1
```

```
Fit method: svdpc
```

```
Number of components considered: 100
```

```
VALIDATION: RMSEP
```

```
Cross-validated using 10 random segments.
```

	(Intercept)	1 comps	2 comps	3 comps	4 comps
cv	12.68	11.65	11.75	8.506	4.211
adjcv	12.68	11.65	11.74	8.500	4.198

```
TRAINING: % variance explained
```

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comp
s						
X	98.65	99.59	99.88	99.99	100.00	100.0
0						
Fat	16.79	16.98	56.44	89.97	93.62	95.2
6						

PCR

- Select 3 components

Coefficients in the original variables

```
pcr.fit1=pcr(f, 3,data=train, validation="none")
summary(pcr.fit1)
coef(pcr.fit1)
scores(pcr.fit1)
l=loadings(pcr.fit1)
print(l,cutoff=0)
Yloadings(pcr.fit1)
plot(pcr.fit1)
```

```
> summary(pcr.fit1)
Data: X dimension: 150 100
      Y dimension: 150 1
Fit method: svdpc
Number of components considered: 3
TRAINING: % variance explained
      1 comps  2 comps  3 comps
X      98.65   99.59   99.88
Fat    16.79   16.98   56.44
```

```
> coef(pcr.fit1)
, , 3 comps
```

	Fat
Channel1	-2.61109872
Channel2	-2.56607281
Channel3	-2.52081831
Channel4	-2.47510841
Channel5	-2.42831883
Channel6	-2.37920922
Channel7	-2.32674365
Channel8	-2.27010925
Channel9	-2.20867856
Channel10	-2.14358670

PCR

```
> l=loadings(pcr.fit1)
> print(l,cutoff=0)
```

Loadings:

	Comp 1	Comp 2	Comp 3
Channel1	-0.079	-0.106	0.089
Channel2	-0.080	-0.108	0.088
Channel3	-0.080	-0.110	0.086
Channel4	-0.081	-0.112	0.084
Channel5	-0.081	-0.113	0.083
Channel6	-0.082	-0.115	0.081
Channel7	-0.082	-0.117	0.079
Channel8	-0.083	-0.118	0.077
Channel9	-0.083	-0.119	0.075
Channel10	-0.084	-0.121	0.073
Channel11	-0.084	-0.122	0.070
Channel12	-0.085	-0.123	0.068

Scores matrix (new coordinates)

```
> scores(pcr.fit1)
```

	Comp 1	Comp 2	Comp 3
155	-9.66855445	0.092805568	-0.1012873027
188	-1.04084544	-0.307978589	-0.3180672681
163	-1.92212872	-0.243714308	0.0124365801
214	1.27768585	-0.232939083	-0.4315433137
97	2.55797242	-0.741279740	0.1582517775
35	-11.17415228	2.124582502	0.2004058835
68	2.96005563	-0.312953959	0.1945895756
106	3.71331162	0.015766621	-0.1130155332

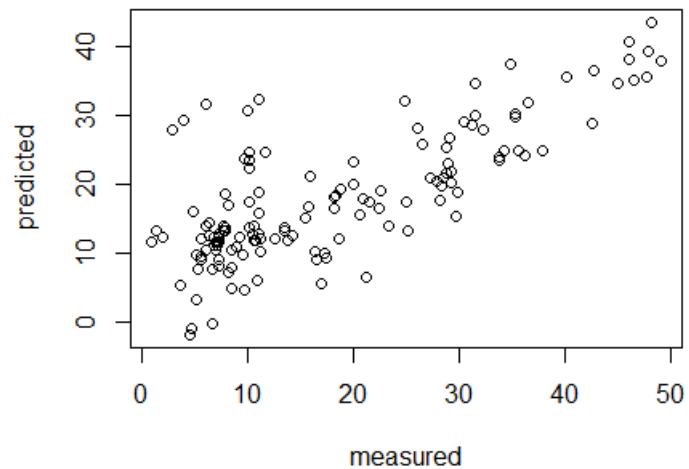
PCR

```
> Yloadings(pcr.fit1)
```

Loadings:

	Comp 1	Comp 2	Comp 3
Fat	-1.015	1.114	-28.847

Fat, 3 comps, train



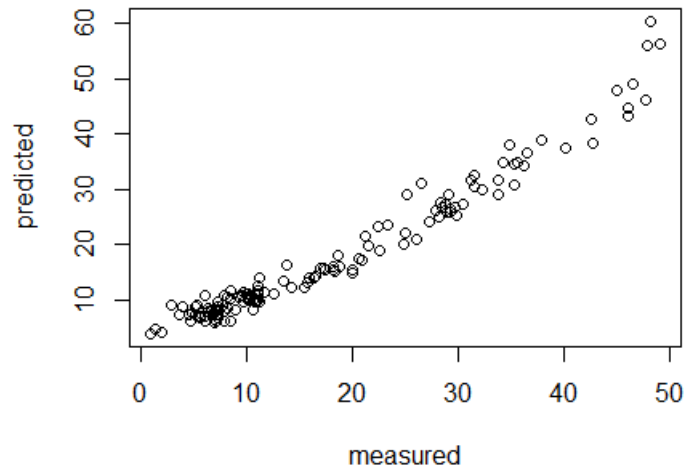
What is the regression equation
in the new coordinates?

Is fit OK?

PCR

- Now 6 components

Fat, 6 comps, train



```
> summary(pcr.fit2)
```

```
Data: X dimension: 150 100
```

```
Y dimension: 150 1
```

```
Fit method: svdpc
```

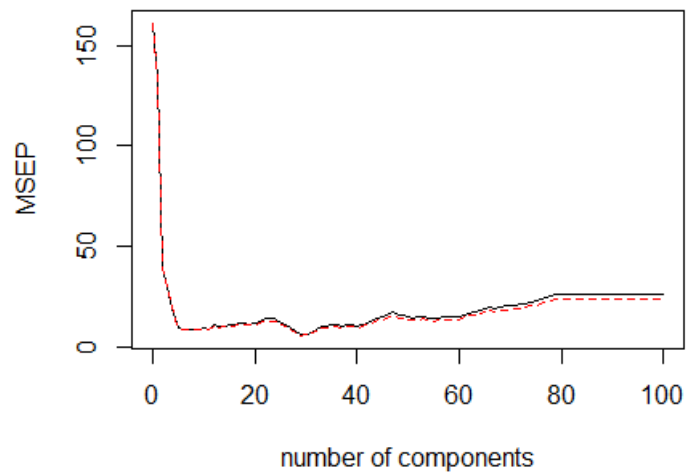
```
Number of components considered: 6
```

```
TRAINING: % variance explained
```

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
X	98.65	99.59	99.88	99.99	100.00	100.00
Fat	16.79	16.98	56.44	89.97	93.62	95.29

PLS

Fat



```
> summary(pls.fit2)
```

```
Data: X dimension: 150 100
```

```
Y dimension: 150 1
```

```
Fit method: kernelpls
```

```
Number of components considered: 6
```

```
TRAINING: % variance explained
```

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
X	98.65	98.95	99.75	99.99	100.00	100.00
Fat	17.10	77.67	83.32	90.58	94.93	95.46

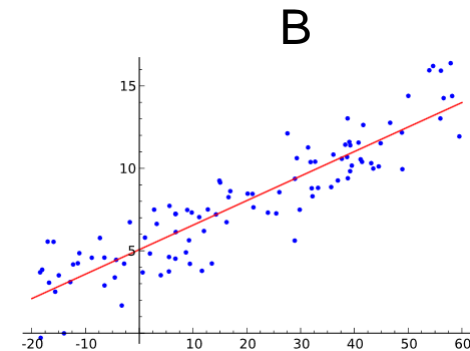
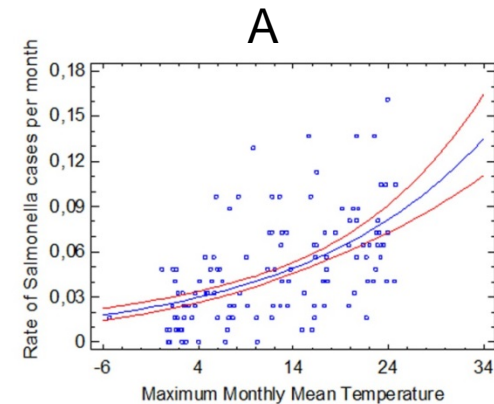
Probabilistic models

- Why it is beneficial to assume a **probabilistic** model?
- A common approach to modelling in CS and engineering:
$$y = f(x, w)$$
- f is known, w is unknown
- Fit model to data with least squares, optimization or ad hoc → find w

Probabilistic models

Arguments against deterministic models:

- The model does not really describe actual data (error is not explained)
 - No difference between modelling data A (Poisson) and B (Normal)
 - Estimation strategy for A is not good for B
- The model typically gives a **deterministic answer**, no information about uncertainty
 - "...The exchange rate tomorrow will be 8.22 ..." 😬



Probabilistic models

Probabilistic model

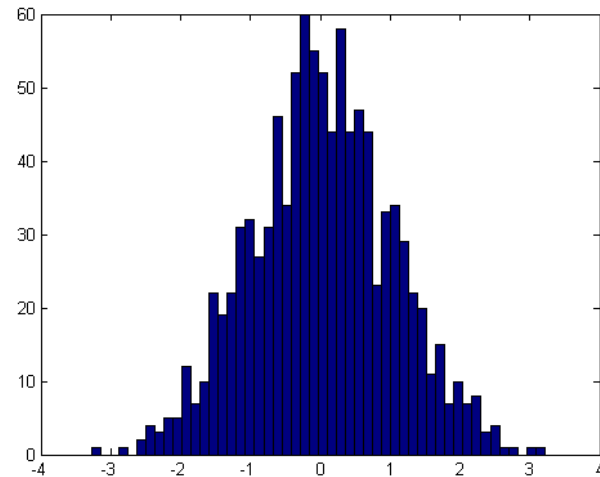
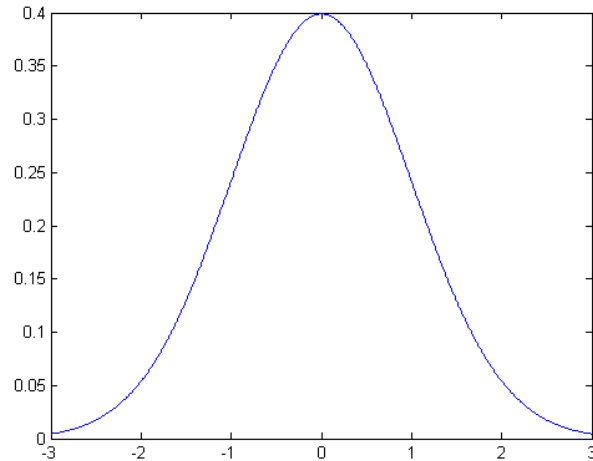
$$Y \sim \text{Distribution}(f(x, w), \theta)$$

- Data is fully explained (error as well)
- Automatic principle for finding parameters: MLE , MAP or Bayes theorem
- Automatic principle for finding uncertainty (conf. limits)
 - **Bootstrap**
 - Posterior probability
- Possibility to generate new data of the same type
 - Further testing of the model

Uncertainty estimation

- Given estimator $\hat{f} = \hat{f}(x, D)$ (or $\hat{\alpha} = \delta(D)$), how to estimate the uncertainty?
- **Answer 1:** if the distribution for data D is given, compute analytically the distribution for the estimator \rightarrow derive confidence limits
 - Often difficult
 - **Example:** In simple linear regression, $\hat{\alpha}$ follows t distribution
- **Answer 2:** Use **bootstrap**

The bootstrap: general principle



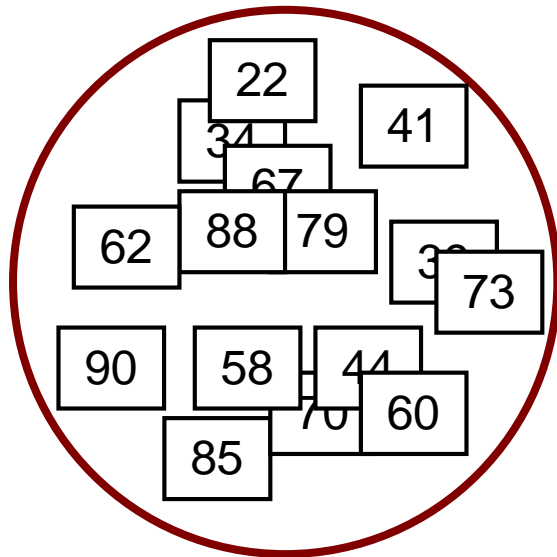
We want to determine uncertainty of $\hat{f}(D, X)$

1. Generate many different D_i from their distribution
2. Use histogram of $\hat{f}(D_i, X)$ to determine confidence limits → unfortunately can not be done (*distr of D is often unknown*)

Instead: Generate many different D_i^* from the empirical distribution (histogram)

Nonparametric bootstrap

Observed data

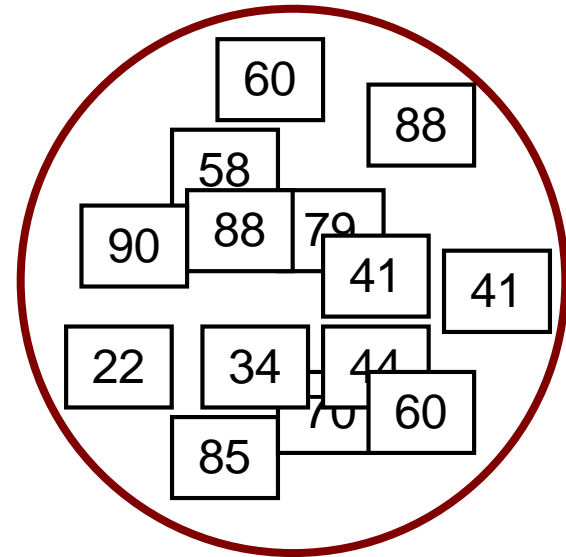


\bar{x}

Sampling with
replacement



Resampled data



$\bar{x}_1^*, \bar{x}_2^*, \dots, \bar{x}_N^*$

Nonparametric bootstrap

Given estimator $\hat{w} = \hat{f}(D)$

Assume $X \sim F(X, w)$, F and w are unknown

1. Estimate \hat{w} from data $\mathbf{D}=(X_1, \dots, X_n)$
2. Generate $\mathbf{D}_1=(X_1^*, \dots, X_n^*)$ by sampling with replacement
3. Repeat step 2 B times
4. The distribution of w is given by $\hat{f}(D_1), \dots, \hat{f}(D_B)$

Nonparametric bootstrap can be applied to any deterministic estimator, distribution-free

Parametric bootstrap

Given estimator $\hat{w} = \hat{f}(D)$

Assume $X \sim F(X, w)$, F is known and w is unknown

1. Estimate \hat{w} from data $\mathbf{D}=(X_1, \dots, X_n)$
2. Generate $\mathbf{D}_1=(X_1^*, \dots, X_n^*)$ by generating from $F(X, \hat{w})$
3. Repeat step 2 B times
4. The distribution of w is given by $\hat{f}(D_1), \dots, \hat{f}(D_B)$

Parametric bootstrap is **more** precise if the distribution form is correct

Uncertainty estimation

1. Get D_1, \dots, D_B by bootstrap
 2. Use $\hat{f}(D_1), \dots, \hat{f}(D_B)$ to estimate the uncertainty
 - Bootstrap percentile
 - Bootstrap Bca
 - ...
- Bootstrap works for all distribution types
 - Can be bad accuracy for small data sets $n < 40$ (empirical is far from true)
 - Parametric bootstrap works even for small samples

Bootstrap confidence intervals

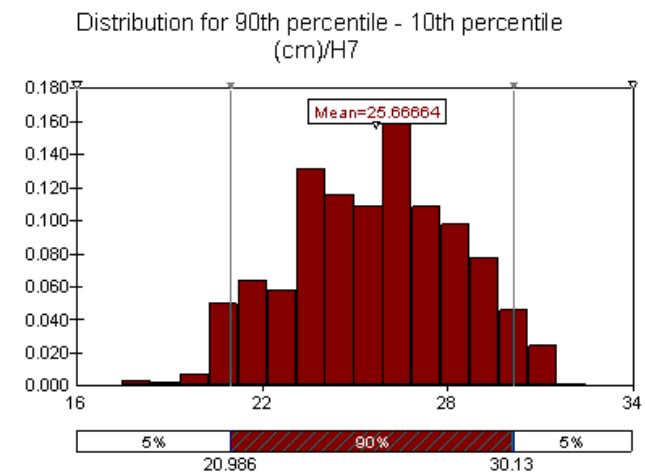
- To estimate $100(1-\alpha)$ confidence interval for w

Bootstrap percentile method

- Using bootstrap, compute $\hat{f}(D_1), \dots, \hat{f}(D_B)$, sort in ascending order, get $w_1 \dots w_B$
- Define $A_1 = \text{ceil}(B \alpha/2)$, $A_2 = \text{floor}(B - B \alpha/2)$
- Confidence interval is given by

$$(w_{A_1}, w_{A_2})$$

Look at the plot...



Bootstrap: regression context

- Model $Y \sim F(X, w)$
- Data $D = \{(Y_i, X_i), i = 1, \dots, n\}$
- Idea: produce several bootstrap sets that are similar to D

Nonparametric bootstrap:

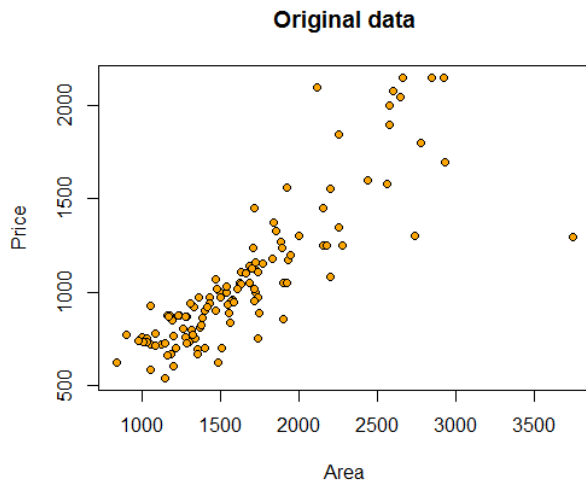
1. Using observation set D , sample **pairs** (X_i, Y_i) with replacement and get bootstrap sample D_1
2. Repeat step 1 B times \rightarrow get D_1, \dots, D_B

Uncertainty estimation

Example: Albuquerque dataset:

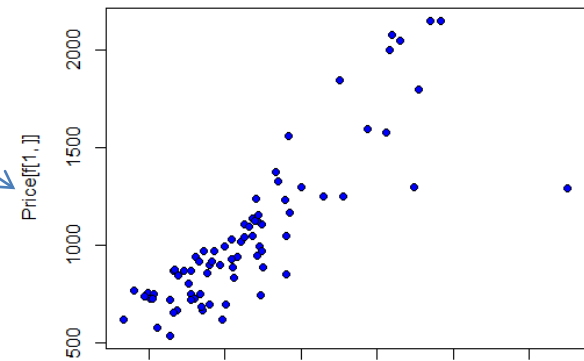
Y=Price of House

X=Area (sqft)

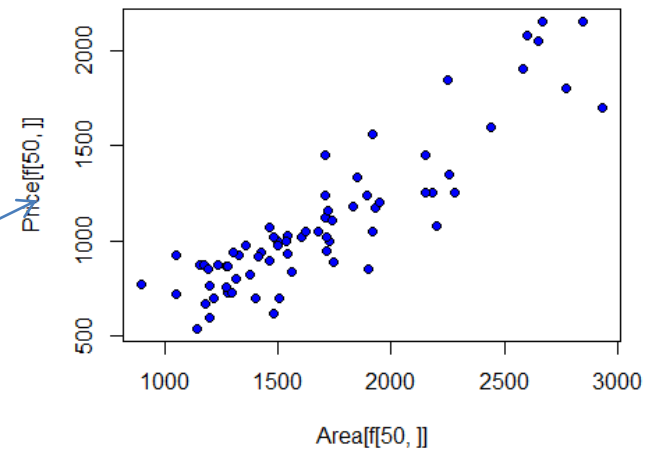


We sample data index,
from $\{1 \dots N\}$

D_1



D_{50}



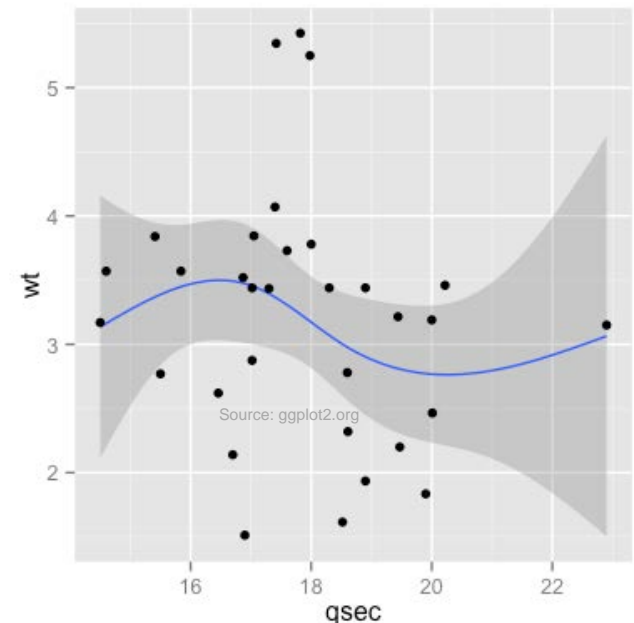
Bootstrap: regression context

Parametric bootstrap

1. Fit a model to $D \rightarrow$ get $\hat{w}(D)$.
2. Set $X_i^* = X_i$, generate $Y_i^* \sim F(X_i, \hat{w})$.
3. $D_i = \{(X_i^*, Y_i^*), i = 1, \dots, n\}$
4. Repeat step 2 B times

Confidence intervals in regression

- Given $Y \sim \text{Distribution}(y|x, w)$, $EY|X = \mu|x = f(x, w)$
 - **Example:** $Y \sim N(w^T x, \sigma^2)$, $\mu|x = f(x, w) = w^T x$
- Estimate intervals for $\mu|x = f(x, w)$ for many X , combine in a **confidence band**
- What is estimator?
 - $\mu|x = f(x, w)$



Confidence intervals in regression

Estimation

1. Compute D_1, \dots, D_B using a bootstrap
2. Fit model to $D_1, \dots, D_B \rightarrow$ estimate $\hat{w}_1, \dots, \hat{w}_B$
3. For a given X , compute $f(X, \hat{w}_1), \dots, f(X, \hat{w}_B)$ and estimate confidence interval by (percentile method)
4. Combine confidence intervals in a band

Bootstrap: R

- Package **boot**

- **Functions:**

- boot()
 - boot.ci() – 1 parameter
 - envelope() – many parameters

- Random random generation for parametric bootstrap:

- Rnorm()
 - Runif()
 - ...

`boot(data, statistic, R, sim = "ordinary",
ran.gen = function(d, p) d, mle = NULL,...)`

Nonparametric bootstrap:

- Write a function *statistic* that depends on *dataframe* and *index* and returns the estimator

```
library(boot)
data2=data[order(data$Area),]#reordering data according
to Area

# computing bootstrap samples
f=function(data, ind){
  data1=data[ind,]# extract bootstrap sample
  res=lm(Price~Area, data=data1) #fit linear model
  #predict values for all Area values from the original
  data
  priceP=predict(res,newdata=data2)
  return(priceP)
}
res=boot(data2, f, R=1000) #make bootstrap
```


Bootstrap: R

Parametric bootstrap:

- Compute value *mle* that estimates model parameters from the data
- Write function *ran.gen* that depends on *data* and *mle* and which generates new data
- Write function *statistic* that depend on *data* which will be generated by *ran.gen* and should return the estimator

```
mle=lm(Price~Area, data=data2)
```

```
rng=function(data, mle) {  
  data1=data.frame(Price=data$Price,  
    Area=data$Area)  
  n=length(data$Price)  
  #generate new Price  
  data1$Price=rnorm(n,predict(mle,  
    newdata=data1),sd(mle$residuals))  
  return(data1)  
}
```

```
f1=function(data1){  
  res=lm(Price~Area, data=data1) #fit linear  
  model  
  #predict values for all Area values from  
  the original data  
  priceP=predict(res,newdata=data2)  
  return(priceP)  
}
```

```
res=boot(data2, statistic=f1, R=1000,  
  mle=mle,ran.gen=rng, sim="parametric")
```

Uncertainty estimation: R

- Bootstrap confidence bands for linear model

```
e=envelope(res) #compute confidence bands
```

```
fit=lm(Price~Area, data=data2)
```

```
priceP=predict(fit)
```

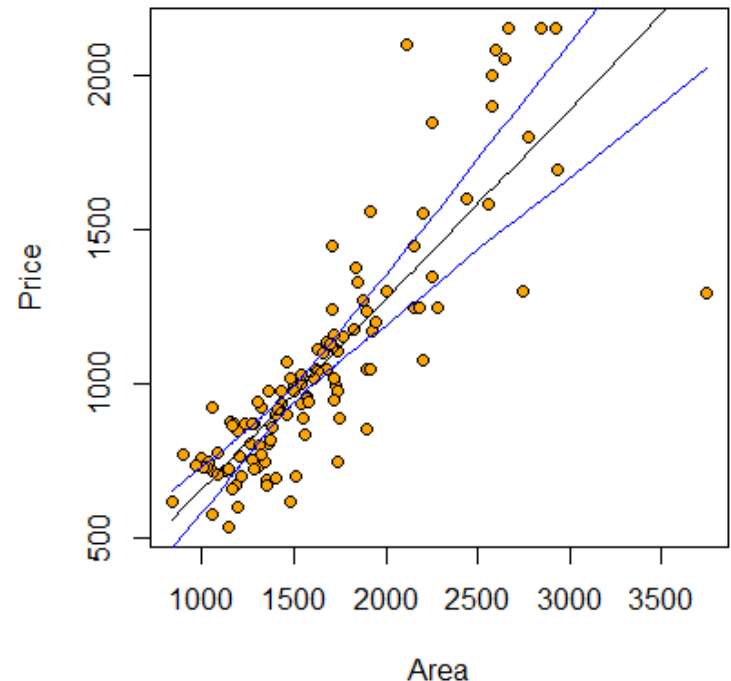
```
plot(Area, Price, pch=21, bg="orange")
```

```
points(data2$Area,priceP,type="l") #plot fitted line
```


```
#plot confidence bands
```

```
points(data2$Area,e$point[2,], type="l", col="blue")
```

```
points(data2$Area,e$point[1,], type="l", col="blue")
```



Prediction bands

- Confidence interval for $Y|X$ = interval for mean $EY|X$
- Prediction interval for $Y|X$ = interval for $Y|X$ 

$$Y \sim \text{Distribution}(x, w)$$

Prediction band for parametric bootstrap

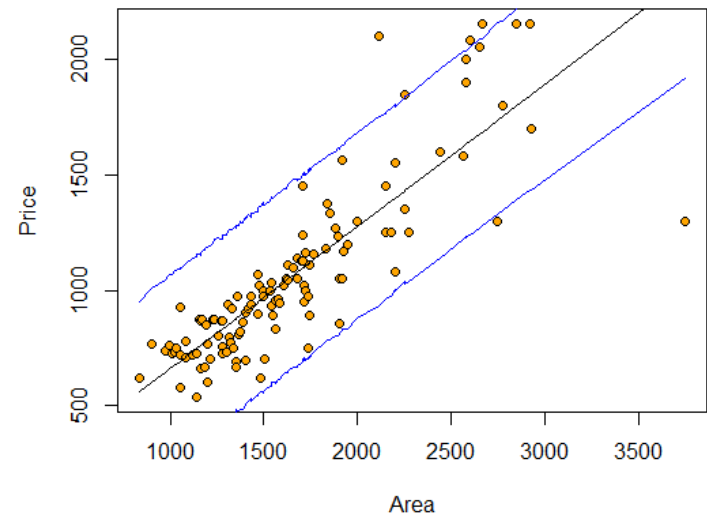
1. Run parametric bootstrap and get D_1, \dots, D_B
2. Fit the model to the data and get $\hat{w}(D_1), \dots, \hat{w}(D_B)$
3. For each X , generate from $\text{Distribution}(X, \hat{w}(D_1)), \dots, \text{Distribution}(X, \hat{w}(D_B))$ and apply percentile method
4. Connect the intervals \rightarrow get the band

Estimation of the model quality

Example: parametric bootstrap

```
mle=lm(Price~Area, data=data2)

f1=function(data1){
  res=lm(Price~Area, data=data1) #fit linear
  model
  #predict values for all Area values from the
  original data
  priceP=predict(res,newdata=data2)
  n=length(data2$Price)
  predictedP=rnorm(n,priceP, sd(mle$residuals))
  return(predictedP)
}
res=boot(data2, statistic=f1, R=10000,
mle=mle,ran.gen=rng, sim="parametric")
```



Why wider band?